


AVL Trees



Performance of BST

Operation	Running Time
$k \text{ in } T$	$O(h)$
$T[k], T[k] = v$	$O(h)$
$T.\text{delete}(p), \text{del } T[k]$	$O(h)$
$T.\text{find_position}(k)$	$O(h)$
$T.\text{first}(), T.\text{last}(), T.\text{find_min}(), T.\text{find_max}()$	$O(h)$
$T.\text{before}(p), T.\text{after}(p)$	$O(h)$
$T.\text{find_lt}(k), T.\text{find_le}(k), T.\text{find_gt}(k), T.\text{find_ge}(k)$	$O(h)$
$T.\text{find_range}(\text{start}, \text{stop})$	$O(s + h)$
$\text{iter}(T), \text{reversed}(T)$	$O(n)$

- Space usage is $O(n)$, where n is the number of items stored in the map.

Binary Search Tree - Best Time

- All BST operations are $O(h)$, where h is tree height.
- In the best case, T has height
$$h = \text{Ceil}(\log(n+1)) - 1$$
- So, best case running time of BST operations is $O(\log n)$


Binary Search Tree - Worst Time

- Worst case running time is $O(N)$
- Insertion of elements in monotonous order

Insert: 2, 4, 6, 8, 10, 12 into an empty BST

- Problem: Lack of **balance**
compare depths of left and right subtree
- Unbalanced degenerate tree

Balanced vs Unbalanced BST



Approaches to balancing trees

- ◆ Don't balance

May end up with some nodes very deep

- ◆ Strict balance

The tree must always be balanced perfectly

- ◆ Pretty good balance

Only allow a little out of balance

- ◆ Adjust on access

Self-adjusting


Balancing Binary Search Trees

Many algorithms exist for keeping BSTs balanced

- Adelson-Velskii and Landis (AVL) trees
(height-balanced trees)
- Splay trees and other self-adjusting trees
- Red-Black trees
- B-trees and other multiway search trees


Perfect Balance

- ◆ Want a **complete tree** after every operation
 - tree is full except possibly in the lower right
- ◆ This is expensive
 - For example, insert 2 in the tree on the left and then rebuild as a complete tree



AVL Tree Definition


- ◆ AVL trees are balanced
- ◆ An AVL Tree is a **binary search tree** such that for every internal node v of T , the **heights of the children of v can differ by at most 1**



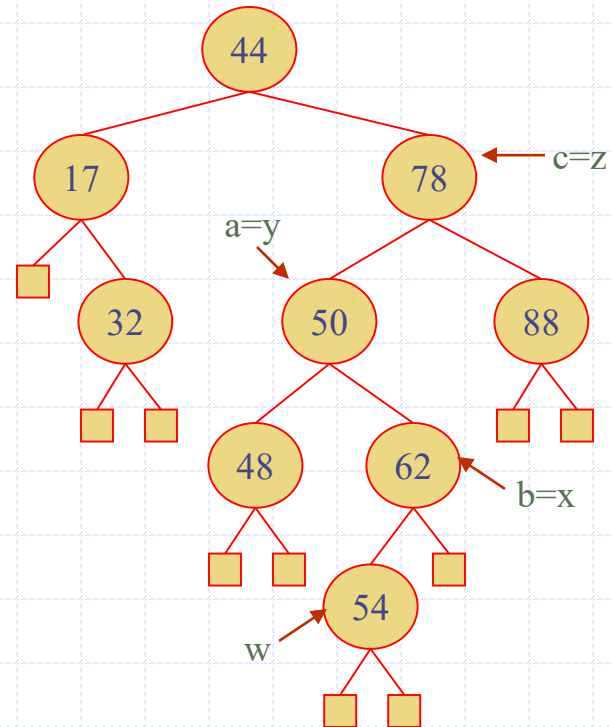
An example of an AVL tree where the heights are shown next to the nodes:

Insertion

- ◆ Insertion is as in a binary search tree
- ◆ Always done by expanding an external node.
- ◆ Example:




before insertion



after insertion

Rotation

Rotate(x): x will become parent, y will become child




Rotate(y): y will become parent, x will become child

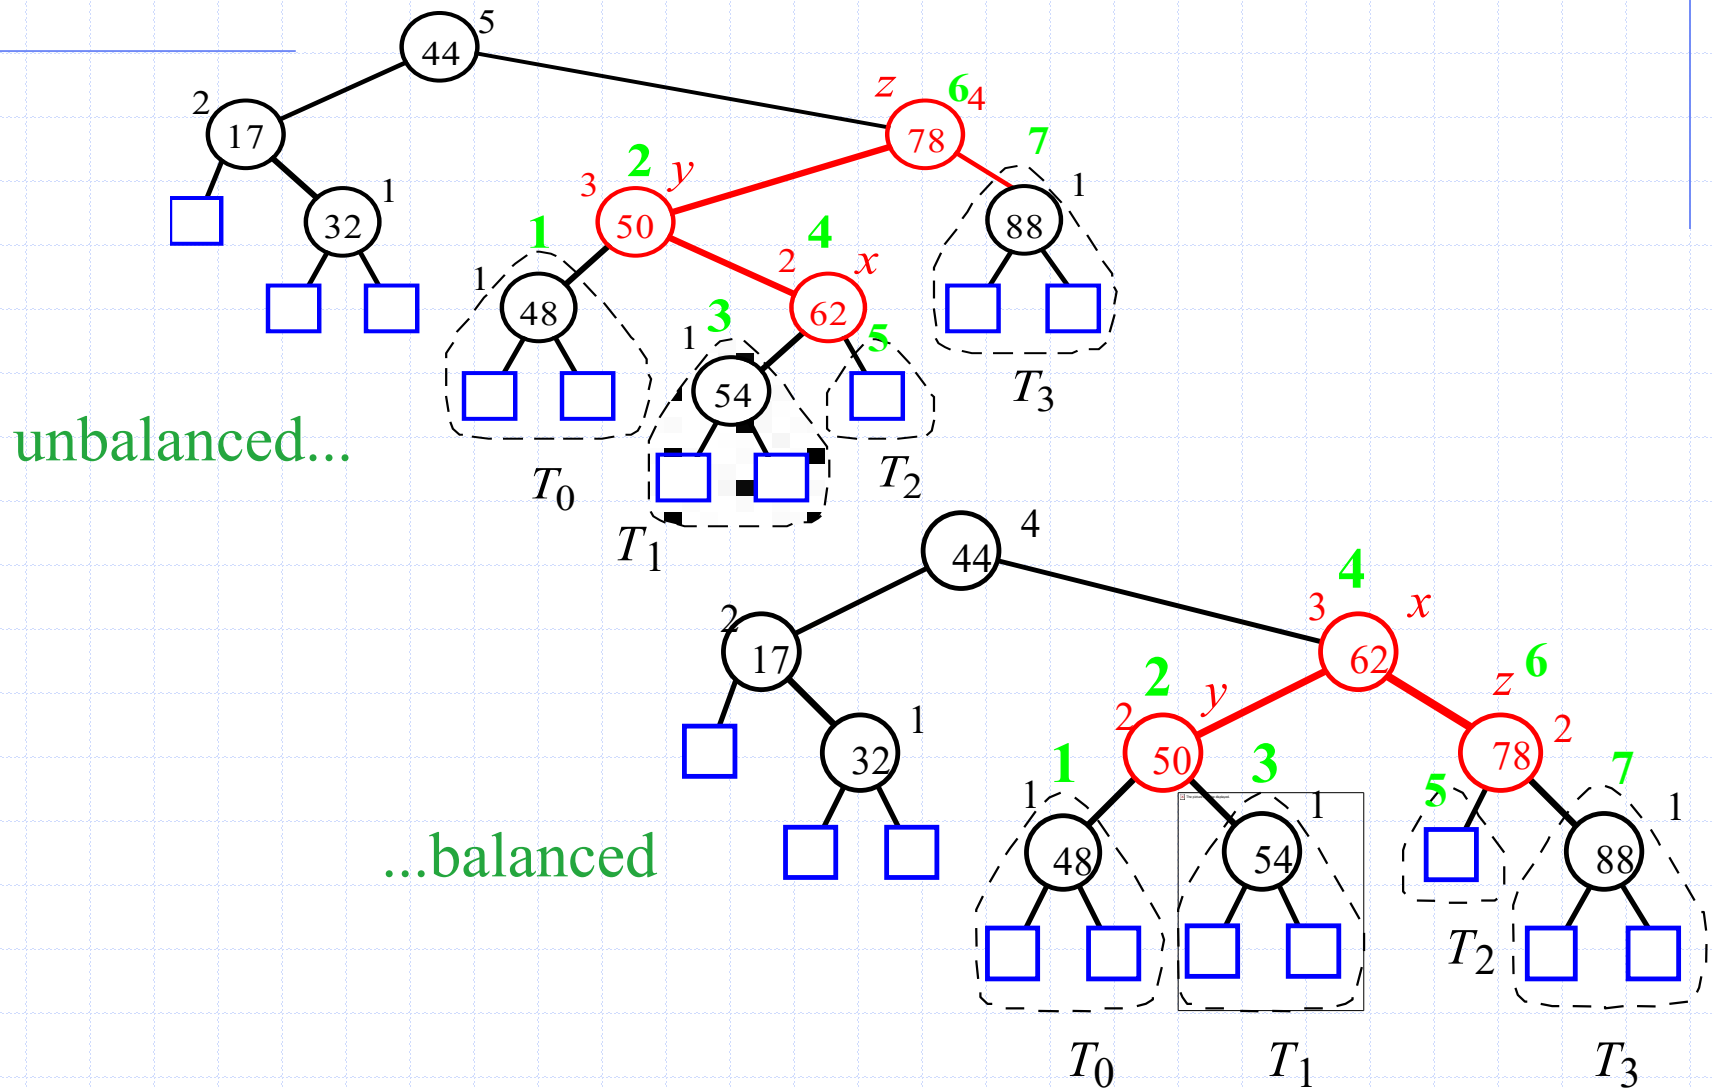
Trinode Restructuring

- ◆ let (a, b, c) be an inorder listing of x, y, z
- ◆ perform the rotations needed to make b the topmost node of the tree

(other two cases are symmetrical)




Insertion Example, continued




Restructuring (as Single Rotations)

◆ Single Rotations: $a < b < c \Rightarrow \text{rotate } (b)$




Restructuring (as Double Rotations)

◆ double rotations: $a < b < c \Rightarrow \text{rotate}(b)$ twice




AVL Insertion: Outside Case

Consider a valid
AVL subtree




AVL Insertion: Outside Case




Inserting into X
destroys the AVL
property at node j

AVL Insertion: Outside Case




Do a “right rotation”:
Node k is promoted &
Node j is demoted

Single Right Rotation



Outside Case Completed




“Right rotation” done!
 (“Left rotation” is mirror
 symmetric)

AVL property has been restored!


AVL Insertion: Inside Case

Consider a valid
AVL subtree




AVL Insertion: Inside Case

Inserting into Y
destroys the
AVL property
at node j



Does “right rotation”
restore balance?


AVL Insertion: Inside Case



“Right rotation”
does not restore
balance... now k is
out of balance


AVL Insertion: Inside Case

Consider the structure
of subtree Y...




AVL Insertion: Inside Case

Y = node i and
subtrees V and W




AVL Insertion: Inside Case




We will do a left-right
“double rotation” . . .

Double Rotation: First Rotation




Double Rotation: Second Rotation



Now do a right rotation



Double Rotation: Second Rotation

right rotation complete
Balance has been restored





Inserted 12 (Single rotation)

Node 70 is imbalanced



From 70, make two steps deeper (choose the deeper path) and make the tri-nodes in-order: 25, 45, 70 with a straight line. So rotate (promote) node 45 once

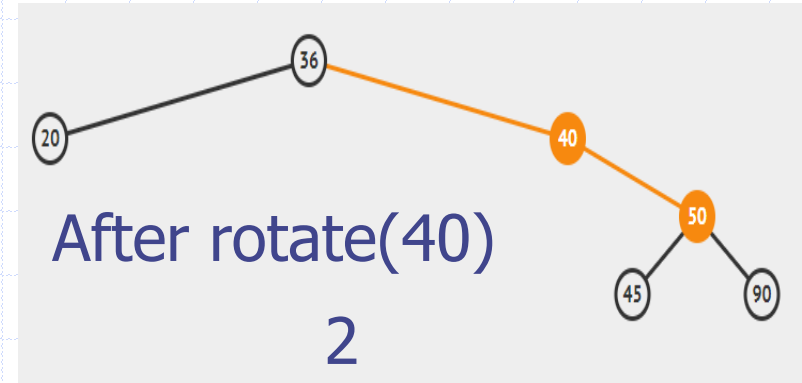



After rotate(45)


Inserted 45 (Double Rotation)

Rotate right, then left

Node 36 is imbalanced

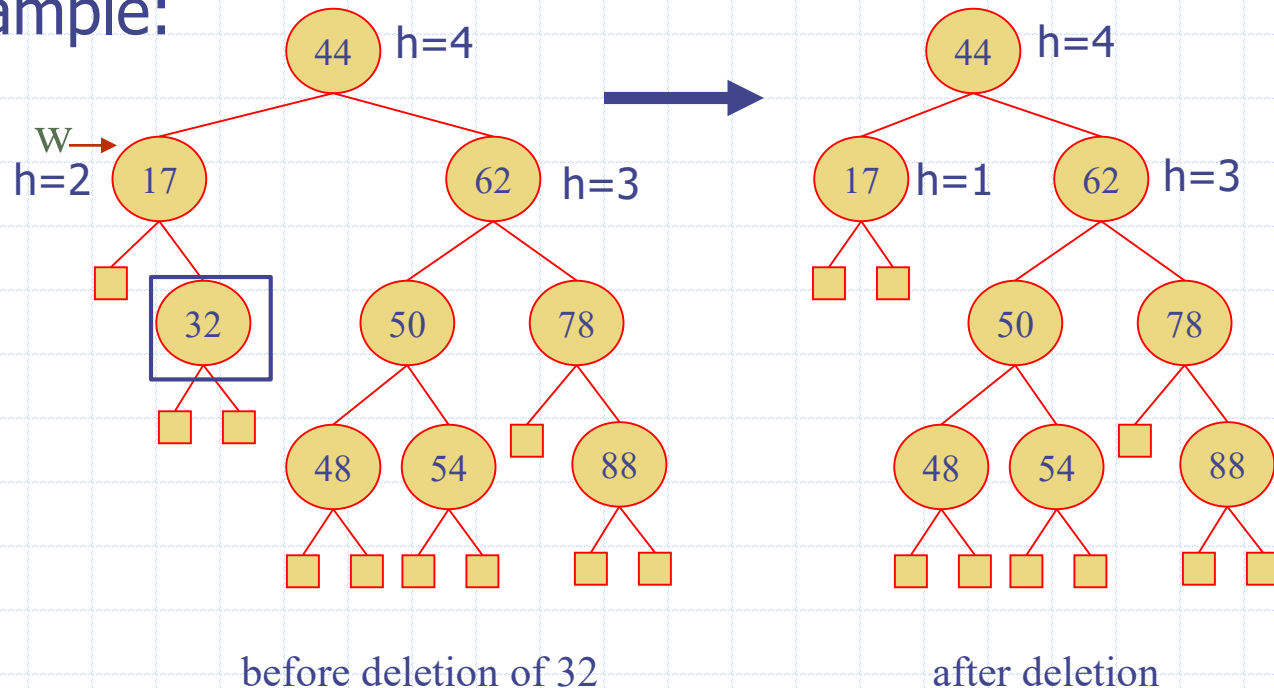


From 36, make two steps deeper (choose the deeper path) and make the tri-nodes in-order: 36, 40, 50 with a non-straight line. So rotate (promote) node 40 twice.




Removal (Delete 32)

- ◆ Removal begins as in a binary search tree, which means the node removed will become an empty external node. Its parent, w , may cause an imbalance.
- ◆ Example:



Rebalancing after a Removal


- ◆ Let z be the first unbalanced node encountered while travelling up the tree from w . Also, let y be the child of z with the larger height, and let x be the child of y with the larger height
- ◆ We perform `restructure(x)` to restore balance at z
- ◆ As this restructuring may upset the balance of another node higher in the tree, we must continue checking for balance until the root of T is reached



Summary

- ◆ Find the unbalanced node from bottom up, say rooted at node z
- ◆ Take the deeper path from z with 2 steps, forming a path $z \rightarrow y \rightarrow x$
- ◆ If $z \rightarrow y \rightarrow x$ is a straight line, rotate(y) once
- ◆ Otherwise, rotate(x) twice

AVL Tree Performance



- ◆ A single restructure takes $O(1)$ time
 - using a linked-structure binary tree
- ◆ Searching takes $O(\log n)$ time
 - height of tree is $O(\log n)$, no restructures needed
- ◆ Insertion takes $O(\log n)$ time
 - initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$
- ◆ Removal takes $O(\log n)$ time
 - initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$

Implementing _Rotate function in Python

```
def _rotate(self, p):
```

```
    """Rotate Position p above its parent.
```

```
    Switches between these configurations, depending on whether p==a or p==b.
```

```
      b
     /\
    a  t2
   /\
  t0 t1
```

```
      a
     /\
    t0 b
   /\
  t1 t2
```

```
    Caller should ensure that p is not the root.
```

```
    """
```

```
    """Rotate Position p above its parent."""
```

```
    x = p._node
```

```
    y = x._parent
```

```
    z = y._parent
```

```
    if z is None:
```

```
        self._root = x
```

```
        x._parent = None
```

```
    else:
```

```
        self._relink(z, x, y == z._left)
```

```
    # now rotate x and y, including transfer of middle subtree
```

```
    if x == y._left:
```

```
        self._relink(y, x._right, True)
```

```
        self._relink(x, y, False)
```

```
    else:
```

```
        self._relink(y, x._left, False)
```

```
        self._relink(x, y, True)
```

```
def _relink(self, parent, child, make_left_child):
```

```
    """Relink parent node with child node (we allow child to be None)."""
```

```
    if make_left_child:
```

```
        parent._left = child
```

```
    else:
```

```
        parent._right = child
```

```
    if child is not None:
```

```
        child._parent = parent
```

```
        # make it a left child
```

```
        # make it a right child
```

```
        # make child point to parent
```

```
        # we assume this exists
```

```
        # grandparent (possibly None)
```

```
        # x becomes root
```

```
        # x becomes a direct child of z
```

```
        # x._right becomes left child of y
```

```
        # y becomes right child of x
```

```
        # x._left becomes right child of y
```


```
        # y becomes left child of x
```

Implementing _restructure function in Python


```
def _restructure(self, x):
    """Perform a trinode restructure among Position x, its parent, and its grandparent.
```

Return the Position that becomes root of the restructured subtree.

Assumes the nodes are in one of the following configurations:



The subtree will be restructured so that the node with key b becomes its root.



Caller should ensure that x has a grandparent.

```
"""
"""Perform trinode restructure of Position x with parent/grandparent."""
y = self.parent(x)
z = self.parent(y)
if (x == self.right(y)) == (y == self.right(z)): # matching alignments
    self._rotate(y)                               # single rotation (of y)
    return y                                       # y is new subtree root
else:                                             # opposite alignments
    self._rotate(x)                               # double rotation (of x)
    self._rotate(x)
    return x                                     # x is new subtree root
```

Implementing _relink function in Python

```
def _relink(self, parent, child, make_left_child):
    """Relink parent node with child node (we allow child to be None)."""
    if make_left_child:
        parent._left = child
    else:
        parent._right = child
    if child is not None:
        child._parent = parent
```

Python Implementation

```
1 class AVLTreeMap(TreeMap):
2     """Sorted map implementation using an AVL tree."""
3
4     #----- nested _Node class -----
5     class _Node(TreeMap._Node):
6         """Node class for AVL maintains height value for balancing."""
7         __slots__ = '_height'          # additional data member to store height
8
9         def __init__(self, element, parent=None, left=None, right=None):
10             super().__init__(element, parent, left, right)
11             self._height = 0           # will be recomputed during balancing
12
13         def left_height(self):
14             return self._left._height if self._left is not None else 0
15
16         def right_height(self):
17             return self._right._height if self._right is not None else 0
```

Python Implementation, Part 2

```
18  #----- positional-based utility methods -----
19  def _recompute_height(self, p):
20      p._node._height = 1 + max(p._node.left_height(), p._node.right_height())
21
22  def _isbalanced(self, p):
23      return abs(p._node.left_height() - p._node.right_height()) <= 1
24
25  def _tall_child(self, p, favorleft=False): # parameter controls tiebreaker
26      if p._node.left_height() + (1 if favorleft else 0) > p._node.right_height():
27          return self.left(p)
28      else:
29          return self.right(p)
30
31  def _tall_grandchild(self, p):    Find a deep path: z -> y -> x
32      child = self._tall_child(p)
33      # if child is on left, favor left grandchild; else favor right grandchild
34      alignment = (child == self.left(p))
35      return self._tall_child(child, alignment)
36
```


Python Implementation, end

```
37 def _rebalance(self, p):
38     while p is not None:
39         old_height = p._node._height      # trivially 0 if new node
40         if not self._isbalanced(p):        # imbalance detected!
41             # perform trinode restructuring, setting p to resulting root,
42             # and recompute new local heights after the restructuring
43             p = self._restructure(self._tall_grandchild(p))
44             self._recompute_height(self.left(p))    Left and right subtree got changed.
45             self._recompute_height(self.right(p))   So recompute heights for them
46             self._recompute_height(p)              # adjust for recent changes
47             if p._node._height == old_height:      # has height changed?
48                 p = None                          # no further changes needed Break the while loop
49             else:
50                 p = self.parent(p)                # repeat with parent
51
52 #----- override balancing hooks -----
53 def _rebalance_insert(self, p):
54     self._rebalance(p)
55
56 def _rebalance_delete(self, p):
57     self._rebalance(p)
```

Continue to maintain balance for parent(p) since height(p) was changed